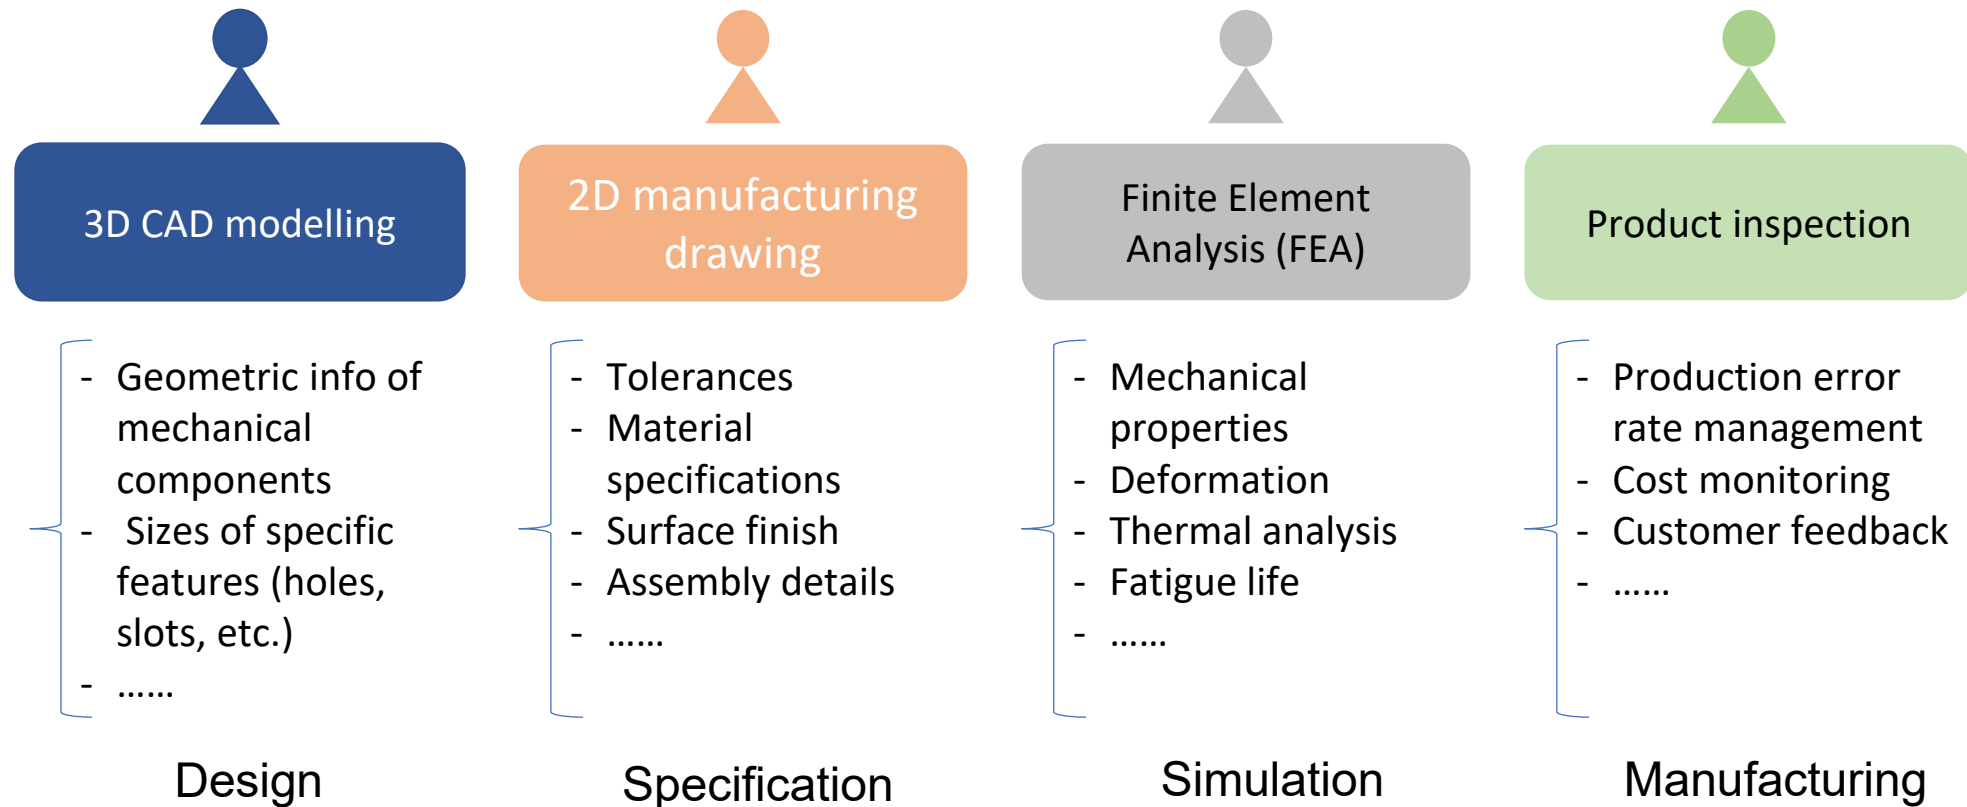


An Agentic LLM with Mechanical Knowledge

Shuang Li, Jonathan Corney*
Erdzan Hodzic, David Eklund#
Rixin Yu, Ralf Denzer^

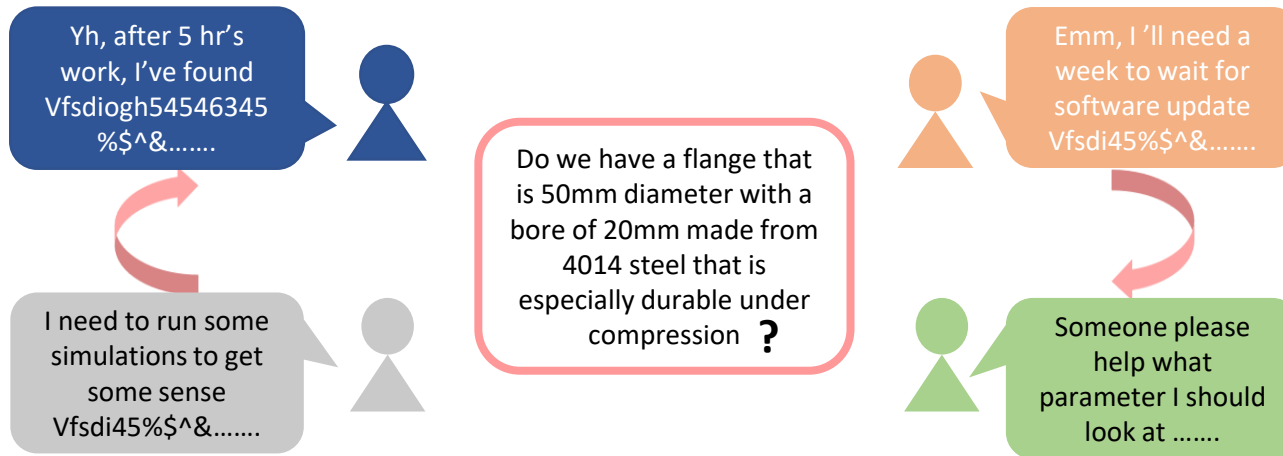
Presented by Dominique Regli
05/2025, MIT

Challenge in Technical Design: CAE Software forms distributed silos of information and knowledge




Challenge in Technical Design: CAE Software forms distributed silos of information and knowledge

Current Practice: The answers to technical questions frequently requires information from multiple silos. Current solutions are facilitated by forms of team chat



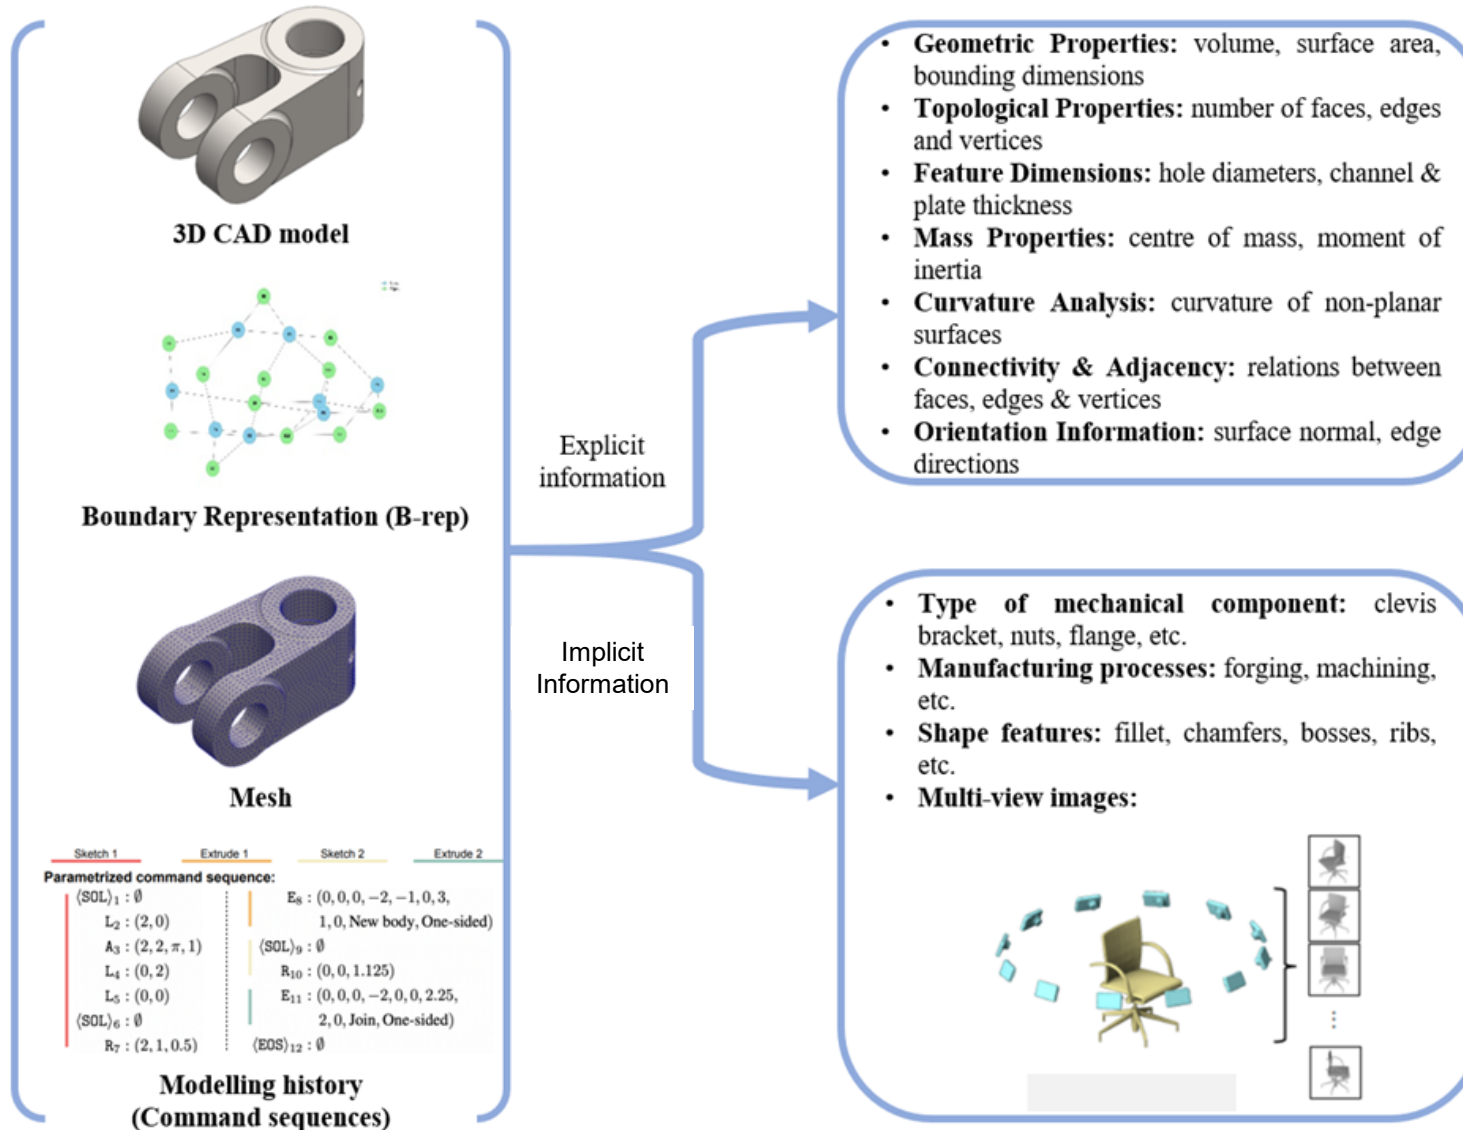
Research Vision: AI responds to questions using holistic knowledge from multiple silos



Classifying Technical Design Questions

Level	Type	Illustrative Examples of MechRAG Prompts	
1	Direct	<p>Query the return values from read or computed from a CAE data files</p> <ul style="list-style-type: none"> What is the volume of part 00210008 in cm³? How many round holes does model 00210337 contain? Does part 00210217 have any type of symmetry? 	Multi-silo conversational interface for data queries
2	Collective	<p>Query combines level 1 information from multiple CAE data structures</p> <ul style="list-style-type: none"> Which parts have the biggest volume and the smallest? What symmetries are present in the dataset of components? 	Multi-silo, multi-design conversational interface for comparative queries
3	Emergent	<p>Query requires response which combines information from multiple sources in ways not explicitly programmed</p> <ul style="list-style-type: none"> Based on part 00217697's modelling history which other components is most similar to it? Which of 0021* series of components would be the most expensive to produce and why? What manufacturing processes is used to produce part 00210215? 	Subjective queries that require product and general knowledge

Tokenising CAE Data



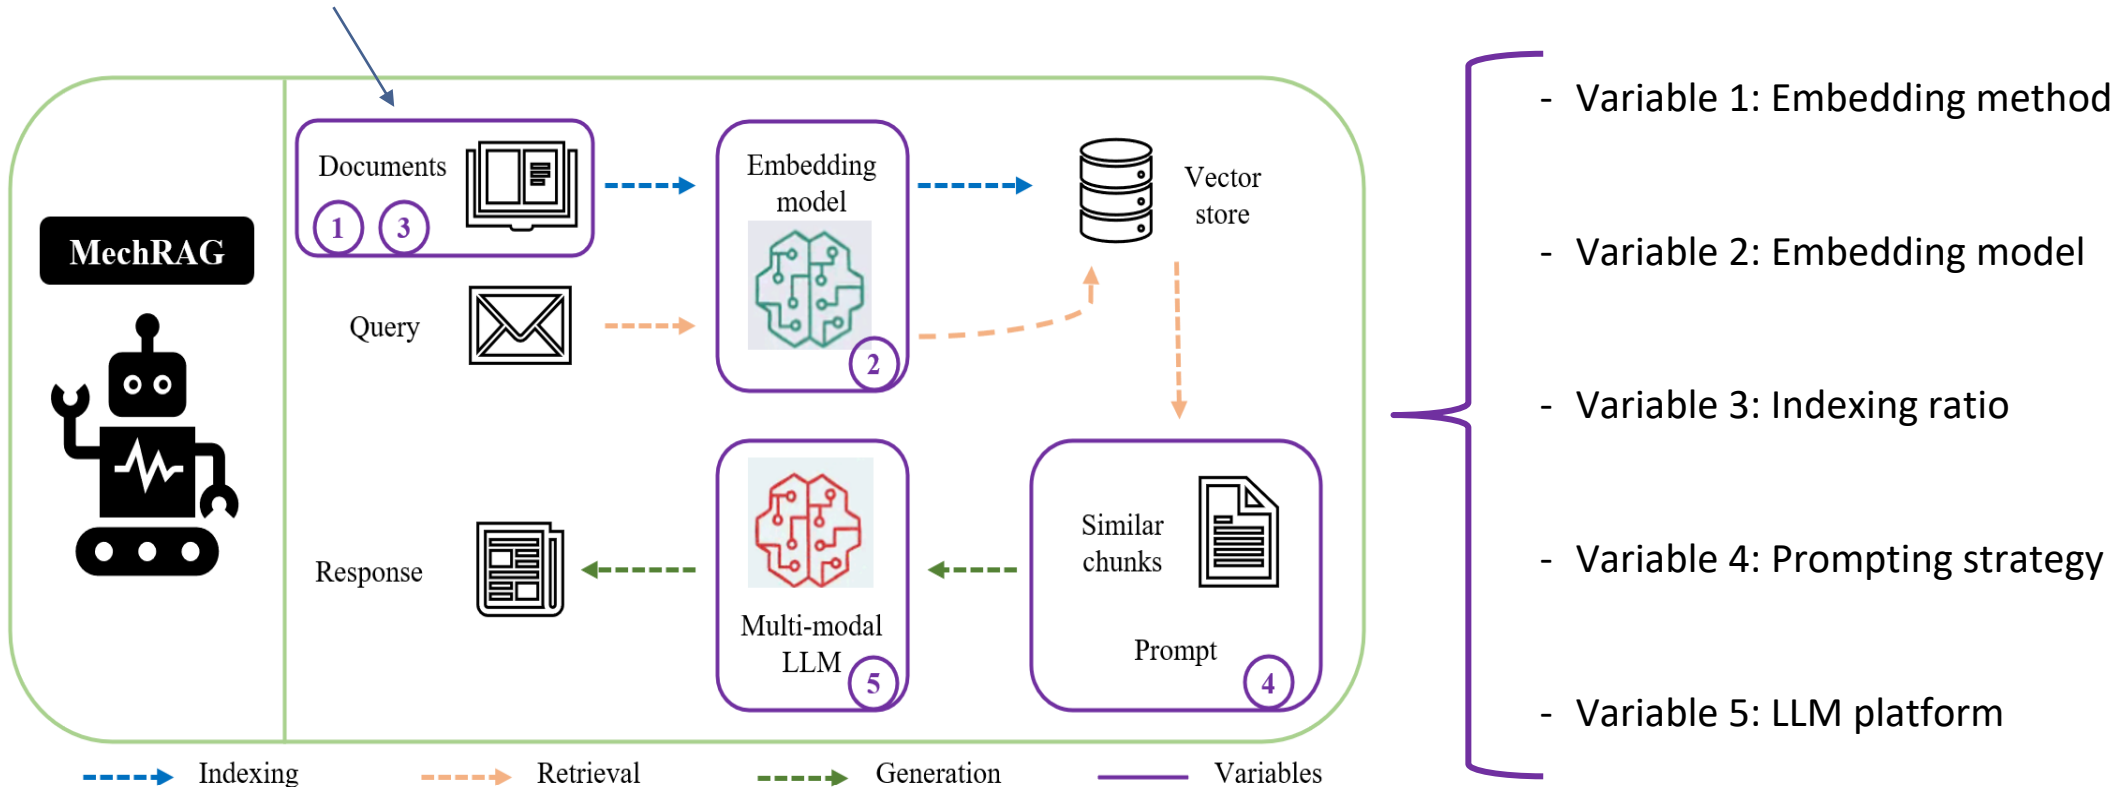
CAE data is typically numerical, so it needs to be processed to provide **summaries** in LLM compatible formats (e.g. images and text).

The volume of the model_00210231 is 237671.829 cubic MM.
 The surface area of the model_00210231 is 38178.907 square MM.
 The volume of this model_00210231 bounding box is 374883.431 cubic MM.
 The variance of the Gaussian curvature of this model_00210231 is 0.008.
 The variance of the Mean curvature of this model_00210231 is 0.000.
 The flat plane of this model_00210231 occupies 0.287%.
 The curved surface of this model_00210231 occupies 99.713%.
 The straight edge of this model_00210231 occupies 24.490%.
 The curved edge of this model_00210231 occupies 75.510%.
 This model_00210231 is radially symmetrical.
 The axis of symmetry of this model_00210231 is [-2.98027006e-10 -1.00000000e
 One of the round holes in this model_00210231 has a diameter 1 of 60.000 MM
 One of the round holes in this model_00210231 has a diameter 2 of 9.000 MM
 One of the round holes in this model_00210231 has a diameter 3 of 9.000 MM
 One of the round holes in this model_00210231 has a diameter 4 of 9.000 MM
 One of the round holes in this model_00210231 has a diameter 5 of 9.000 MM
 One of the round holes in this model_00210231 has a diameter 6 of 60.000 MM

Fragment of a **text summary** that combines information from different CAE representations (e.g. BRep, Mesh, Design History etc).

MechRAG architecture

Multi-modal summary
of CAE data

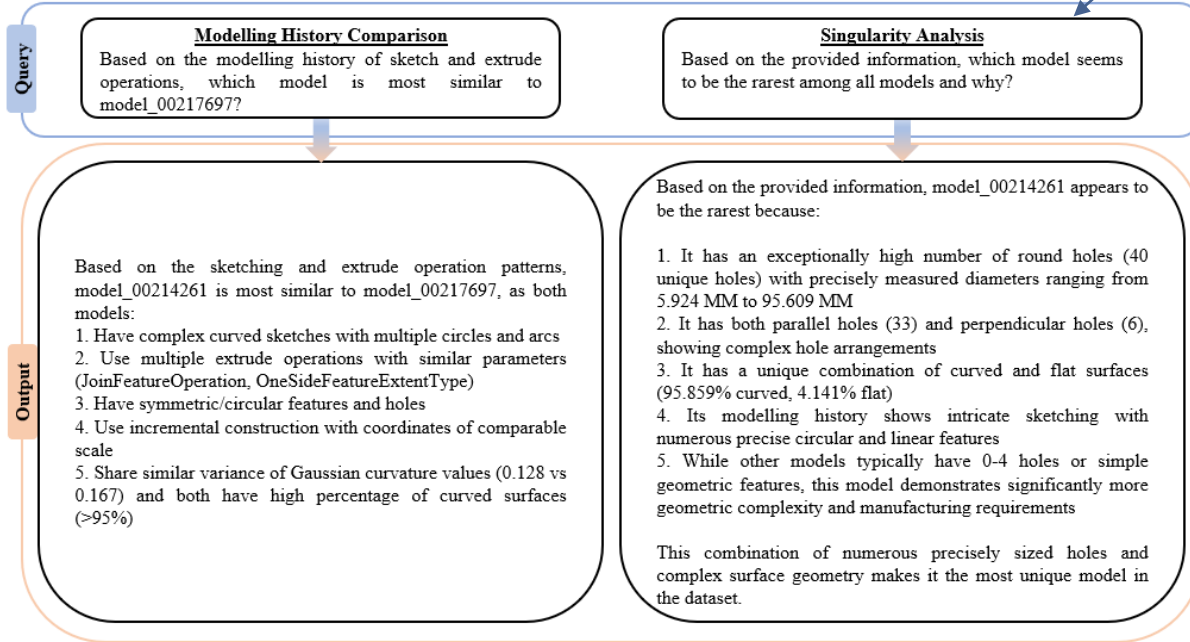


The **Retrieval-Augmented Generation (RAG)** mechanism enhances large language models (LLMs) by combining their generation capabilities with external information retrieval. Instead of relying solely on their internal training data, RAG models first retrieve relevant documents or knowledge snippets from an external database or knowledge base based on the input query. They then use this retrieved information to ground and inform their generated responses, leading to more accurate, up-to-date, and contextually relevant outputs without the need for retraining the LLM itself.


MechRAG example generations

Ambiguous prompts: “most similar” & “rarest”

ChatGPT Prompts




00217697



00214261

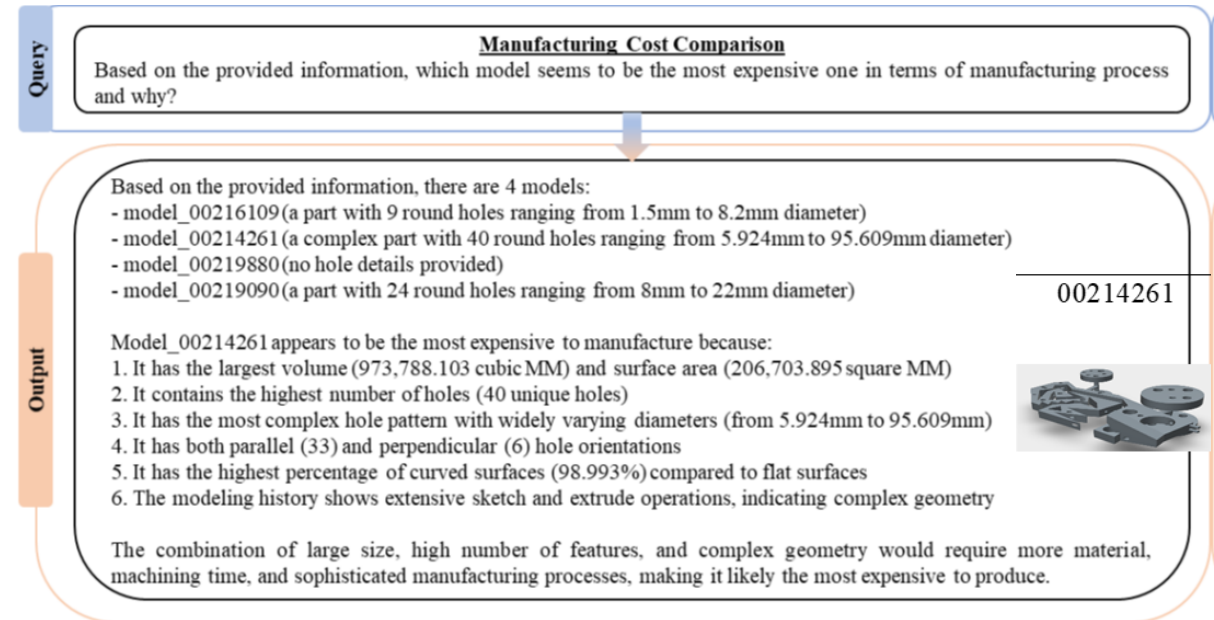


00214261



ChatGPT responds with mix of CAE data (via RAG) and general knowledge

Ambiguous prompts: “most expensive to manufacture”




00214261



Future Work

- ❖ **Token limits of the context window in LLMs:** it's not unusual for engineering companies have tens of thousands of CAD models, but it's possible to exceed an LLMs token limit with less than a hundred models. Consequently work is needed to minimise the number of tokens required to represent engineering product knowledge.



- ❖ **How to utilise other modalities/formats of data** (e.g. FEA/CFD simulation and process) to augment LLMs?

Thanks for your attention!